
Visual Adversarial Attacks and Defenses in the Physical World: A Survey

www.surveyx.cn

Abstract

Visual adversarial attacks and defenses in the physical world pose significant threats to the security and reliability of various applications, including self-driving cars, surveillance systems, and facial recognition technology. This survey provides a comprehensive review of adversarial attack and defense methods, highlighting the importance of establishing robust defenses within computer-aided systems to ensure safety and maintain trust in automated assessments. The study explores the concept of adversarial examples, their generation methods, and the various types of adversarial attacks that can occur in the physical world, including attacks on image classification, object detection, and segmentation. Defense strategies against adversarial attacks, including adversarial training methods and input preprocessing and detection methods, are also discussed. The survey emphasizes the need for continued research in this area to develop more effective defense strategies and ensure the security and reliability of computer vision systems. By understanding the complexities of visual adversarial attacks and defenses, researchers and practitioners can better appreciate the importance of developing robust defense strategies to protect against these attacks and ensure the safety and security of various applications.

1 Introduction

1.1 Background and Context

The importance of visual adversarial attacks and defenses in the physical world is underscored by their significant implications for the security and reliability of various applications, including self-driving cars, surveillance systems, and facial recognition technology, as noted by [1], [2], and [3]. The concept of visual adversarial attacks and defenses is crucial in understanding the importance of adversarial examples in computer vision security, with background information provided in [4]. This highlights the potential impact of these attacks on computer vision security and emphasizes the concept of visual adversarial attacks and defenses in the physical world. Furthermore, the study of visual adversarial attacks and defenses is crucial for ensuring the reliability and security of various applications that rely on computer vision, with [1], [2], and [3] emphasizing the need for robust defense strategies and the development of more secure computer vision systems.

1.2 Significance and Impact

The significance and potential impact of visual adversarial attacks and defenses on computer vision security are substantial, with far-reaching implications for the reliability and trustworthiness of various applications, including self-driving cars, surveillance systems, and facial recognition technology, as discussed in [1], [2], and [3]. According to [4], the potential impact of visual adversarial attacks on computer vision security is significant, and it is essential to investigate and develop effective defense strategies to mitigate these risks. The study of visual adversarial attacks and defenses is crucial for ensuring the security and reliability of computer vision systems, with [1] and [3] highlighting the




Figure 1: chapter structure

importance of understanding the significance of visual adversarial attacks and defenses in the physical world. Additionally, [2] emphasizes the need for developing effective defense strategies to ensure the security and reliability of computer vision systems, with the potential risks and consequences of adversarial attacks on computer vision security being significant, as noted by [1], [2], and [3].

1.3 Structure of the Survey

The survey is organized into six main sections, each focusing on a distinct aspect of visual adversarial attacks and defenses in the physical world, with [1] and [3] providing an overview of the concept of visual adversarial attacks and defenses, including their significance and potential impact on computer vision security. The background and definitions section, informed by [4] and [2], defines key terms and concepts related to visual adversarial attacks and defenses, including adversarial machine learning and computer vision. The section on adversarial attacks in the physical world, examined in [1] and [3], discusses the different types of adversarial attacks that can occur in the physical world, including attacks on image classification, object detection, and segmentation. The defenses against adversarial attacks section, discussed in [2] and [4], presents various defense strategies against adversarial attacks, including adversarial training methods and input preprocessing and detection methods. The case studies and applications section, drawing on [1] and [3], presents case studies and applications of visual adversarial attacks and defenses in real-world scenarios, such as autonomous vehicles, surveillance systems, and facial recognition technology. Finally, the conclusion section, informed by [2] and [4], summarizes the key findings and takeaways from the survey, emphasizing the importance of understanding and addressing visual adversarial attacks and defenses in the physical world, as emphasized by [1] and [3].

1.4 Research Questions and Objectives

The research questions and objectives of this survey are centered around understanding the concept of visual adversarial attacks and defenses in the physical world, as highlighted in [1], [2], and [3]. The primary research question is: What are the key challenges and limitations of existing defense strategies against visual adversarial attacks in the physical world, and how can they be addressed to ensure the security and reliability of computer vision systems, as discussed in [4] and [2]? The survey also aims to investigate the most effective methods for generating adversarial examples in the physical world, and how they can be used to evaluate the robustness of computer vision systems, as explored in [1] and [3]. Additionally, the survey examines the potential risks and consequences of adversarial attacks on computer vision security in real-world applications, such as autonomous vehicles, surveillance systems, and facial recognition technology, as noted by [1], [2], and [3]. The objectives of this survey are to provide a comprehensive review of adversarial attack and defense methods, to identify the key challenges and limitations of existing defense strategies, and to highlight the significance of establishing robust defenses within computer-aided systems to ensure safety and maintain trust in automated assessments, as emphasized by [1], [2], and [3]. The following sections are organized as shown in Figure 1.

2 Background and Definitions

2.1 Introduction to Adversarial Examples

The concept of adversarial examples is crucial in understanding the importance of visual adversarial attacks and defenses in computer vision security, as highlighted in [5], [6], [7], [8], [9], [10], [11], [3], and [4]. Adversarial examples are inputs to a neural network that are specifically designed to cause the network to make a mistake, as defined in [6] and [8]. According to [7] and [9], small, imperceptible perturbations to the input data can cause deep neural networks to produce misclassified predictions with high confidence. The definition of adversarial examples, as discussed in [10] and [5], is critical in understanding their importance in computer vision security. [1] and [3] emphasize the concept of adversarial examples and their importance in computer vision security. Furthermore, [4] provides a comprehensive introduction to adversarial examples, including their definition, significance, and importance in computer vision security. Neural network-based image classifiers are susceptible to adversarial examples, which are minutely perturbed inputs that fool classifiers, as noted in [6]. The concept of adversarial examples is closely related to the importance of robustness in computer vision security, as discussed in [1] and [3], highlighting the concept of adversarial machine learning and its impact on machine learning models in cyberphysical systems. Understanding the concept of adversarial examples and their importance in computer vision security is essential for developing robust defense strategies and more secure computer vision systems.

2.2 Key Terms and Concepts

The study of visual adversarial attacks and defenses relies on a deep understanding of key terms and concepts related to adversarial machine learning and computer vision, as defined in [11], [12], [13], [14], [15], [16], [17], [18], [19], and [20]. Adversarial machine learning refers to the process of designing and constructing machine learning models that are resilient to adversarial attacks, as discussed in [11] and [12]. Computer vision refers to the field of study that focuses on the development of algorithms and statistical models that enable computers to interpret and understand visual data from the world, as defined in [15] and [16]. The concept of adversarial examples is crucial in understanding the importance of visual adversarial attacks and defenses, as highlighted in [13] and [18]. Adversarial examples can be used to evaluate the robustness of computer vision systems, as noted in [19] and [20]. The study of visual adversarial attacks and defenses also relies on an understanding of key terms and concepts related to face recognition, physical-world condition variations, and environmental condition variations, as discussed in [17]. The definitions of key terms and concepts related to visual adversarial attacks and defenses, including adversarial machine learning and computer vision, are essential in understanding the significance of this research area, as emphasized by [11] and [12]. By understanding these key terms and concepts, researchers and practitioners can better appreciate the importance of developing robust defense strategies against visual adversarial attacks and the need for continued research in this field, as highlighted in [14]. The survey organizes the current methods or research into several fields or stages, including construction of adversarial examples,

evaluation metrics, defense techniques, and future research directions, as discussed in [15]. The primary challenge in addressing adversarial examples is the lack of understanding of their cause and the difficulty in constructing effective defense mechanisms, as noted in [18].

2.3 Adversarial Machine Learning and Computer Vision

The relationship between adversarial machine learning and computer vision is crucial in understanding the importance of visual adversarial attacks and defenses in the physical world, as highlighted in [21], [22], [23], [24], [25], [26], [14], [27], [28], and [29]. According to [21], the relationship between adversarial machine learning and computer vision is vital in understanding the importance of visual adversarial attacks and defenses. The impact of adversarial examples on computer vision is significant, as noted in [22], which discusses object detectors such as Faster RCNN and YOLO. The survey [24] organizes the current methods and research into different fields and stages, including the analysis of explanation methods, the design of adversarial attacks, and the identification of the requirements that adversarial attacks should satisfy. Furthermore, [23] categorizes types of adversarial machine learning attacks, and [25] notes that the high curvature of a decision boundary makes adversarial examples decrease their ability to attack other models, especially adversarial robust models. The relationship between adversarial machine learning and computer vision is also discussed in [26], which includes methods for generating adversarial examples. As [14] emphasizes, the survey introduces a novel perspective on adversarial robustness by categorizing data properties into different dimensions, such as problem setup, data property, and practicality. The impact of adversarial machine learning on computer vision security is significant, as highlighted in [27] and [28]. Additionally, [29] introduces a novel perspective on understanding adversarial examples, focusing on the relationship between generalization and robustness, and exploring the impact of various factors on the vulnerability of deep neural networks. The relationship between adversarial machine learning and computer vision can also be understood through the use of geometric proofs to show the relationship between the norm of perturbations and the input dimension, as discussed in [30].

2.4 Methods for Generating Adversarial Examples

The generation of adversarial examples is a crucial aspect of understanding the importance of visual adversarial attacks and defenses in the physical world. Various methods have been proposed to generate adversarial examples, including perturbation-based and generative model-based approaches [31], [32], [33], [34], [35], and [36]. Perturbation-based methods involve adding small perturbations to the input data to create adversarial examples, while generative model-based approaches use generative models to generate new data that can fool the machine learning model. According to [37], adversarial examples can be generated using various methods, such as the Fast Sign method, Iterative methods, and L-BFGS method. The L-BFGS attack, Fast Gradient Sign Method, and Basic Iterative Method are also commonly used methods for generating adversarial examples, as discussed in [38] and [39]. Additionally, the proposed algorithm, Robust Physical Perturbations (RP2), generates perturbations robust to widely changing distances and angles of the viewing camera, by modeling physical dynamics and synthetic transformations [40]. The Expectation Over Transformation (EOT) method is also used to synthesize adversarial examples robust over a chosen distribution of transformations, using a 3D rendering process and stochastic gradient descent [41]. Spatial transformation can be used to generate adversarial examples, and the stAdv method uses spatial transformers to deform the input images and produce more realistic adversarial examples [42]. Furthermore, methods for generating adversarial examples include the proposed method of designing a universal perturbation used in the adversarial T-shirt to deceive a single object detector, and then proposing a min-max (robust) optimization framework to design the universal adversarial patch against multiple object detectors [43]. The Fast Gradient Sign Method and iterative approaches are also used to generate adversarial examples [44]. ABBA uses kernel prediction and motion blur generation to create adversarial examples, which is a method for generating adversarial examples [45]. A novel approach called AdvDrop crafts adversarial examples by dropping existing details from clean images in the frequency domain, rather than adding adversarial perturbation in the spatial domain [35]. Perturbing the disentangled latent representations of an input image to generate adversarial examples, using a common feature variational autoencoder (CF-VAE) and a generative adversarial network-based booster (GAN-B) is also a method for generating adversarial examples [13]. The proposed methods, AdaBelief Iterative Fast Gradient Method (ABI-FGM) and Crop-Invariant attack Method (CIM), are used to generate more transferable adversarial examples [46].

2.5 Importance of Understanding Adversarial Examples

The importance of understanding adversarial examples in the context of real-world applications cannot be overstated, as they have significant implications for the security and reliability of various applications, including self-driving cars, surveillance systems, and facial recognition technology. According to [47], the use of natural adversarial examples to evaluate the robustness of deep learning models has several advantages, including their ability to reveal the weaknesses of current models. This is crucial in understanding the importance of adversarial examples, as it allows researchers and practitioners to identify the limitations of current models and develop more robust defense strategies. Furthermore, [29] highlights the strengths of deep neural networks in terms of their robustness, but also discusses the limitations and weaknesses of current research, including the lack of understanding of the underlying causes of adversarial examples. By understanding adversarial examples and their importance in computer vision security, we can better appreciate the need for robust defense strategies and the development of more secure computer vision systems. The survey aims to provide a comprehensive review of adversarial attack and defense methods, highlighting the significance of establishing robust defenses within computer-aided systems to ensure safety and maintain trust in automated assessments. As [47] emphasizes, the use of natural adversarial examples is essential in evaluating the robustness of deep learning models, and [29] highlights the need for continued research in this area to understand the underlying causes of adversarial examples. The importance of understanding adversarial examples in the context of real-world applications is vital, and continued research in this area is necessary to ensure the security and reliability of computer vision systems.

The concept of adversarial attacks in the physical world is complex and multifaceted. To better understand the organization and implications of these attacks, it is helpful to visualize their hierarchical structure. Figure 2 illustrates this structure, starting with an introduction to adversarial attacks that provides an overview of the concept and its importance in the physical world. The figure then branches out into types of attacks, which include specific examples such as attacks on image classification, object detection, and segmentation. These categories are further explored through case studies that present real-world examples of adversarial attacks. The case studies encompass a range of scenarios, including robust physical perturbations and the deception of computer vision systems, as well as more nuanced examples like adversarial T-shirts and object detectors. By examining the hierarchical structure of adversarial attacks as shown in Figure 2, researchers can gain a deeper understanding of the threats posed by these attacks and the need for robust defenses against them.

3 Adversarial Attacks in the Physical World

3.1 Introduction to Adversarial Attacks in the Physical World

The concept of adversarial attacks in the physical world is crucial in understanding the importance of visual adversarial attacks and defenses in computer vision security, as highlighted in [4]. Adversarial attacks in the physical world refer to the process of designing and constructing physical objects or environments that can manipulate the behavior of computer vision systems, such as self-driving cars, surveillance systems, and facial recognition technology [4]. The physical world provides a unique set of challenges and opportunities for adversarial attacks, as the attacker must consider the physical properties of the object or environment being attacked, as well as the limitations and capabilities of the computer vision system being targeted [4]. Understanding the concept of adversarial attacks in the physical world is essential in developing robust defense strategies against these attacks, as the potential consequences of a successful attack can be severe [4]. By understanding the concept of adversarial attacks in the physical world, researchers and practitioners can better appreciate the importance of developing secure and robust computer vision systems that can withstand these types of attacks [4].

3.2 Types of Adversarial Attacks in the Physical World

The physical world provides a unique set of challenges and opportunities for adversarial attacks, and various types of attacks can occur in this domain, including attacks on image classification, object detection, and segmentation [8, 20, 9, 10, 1, 2, 3]. According to [48, 49, 50, 51, 52, 53, 54, 55], the types of adversarial attacks in the physical world, such as attacks on image classification, object




Figure 2: This figure illustrates the hierarchical structure of adversarial attacks in the physical world, including introduction to adversarial attacks, types of attacks, and case studies. The introduction provides an overview of the concept and importance of adversarial attacks in the physical world, while the types of attacks section delves into specific examples such as attacks on image classification, object detection, and segmentation. The case studies section presents real-world examples of adversarial attacks, including robust physical perturbations and deception of computer vision systems, as well as adversarial T-shirts and object detectors.

detection, and segmentation, can have severe consequences in real-world applications, such as self-driving cars, surveillance systems, and facial recognition technology. The problem of generating transferable adversarial examples that can effectively attack deep neural networks under the black-box setting is also a significant challenge, as noted in [46], where the main difficulty lies in improving the transferability of adversarial examples, which is crucial for successful black-box attacks [25]. Additionally, [6] presents a method for generating targeted adversarial examples in the query-limited, partial-information, and label-only settings, and demonstrates the vulnerability of machine learning systems to black-box attacks. As illustrated in Figure 3, the hierarchical structure of adversarial attacks in the physical world encompasses image classification, object detection, and segmentation, with corresponding attack and defense strategies. This hierarchical structure highlights the complexity and interconnectedness of these attacks, emphasizing the need for comprehensive defense strategies

to ensure the security and reliability of computer vision systems. The various types of adversarial attacks that can occur in the physical world, including attacks on image classification, object detection, and segmentation, are a significant concern, and understanding these attacks is crucial in developing effective defense strategies [8, 20, 9, 10, 1, 2, 3].




Figure 3: This figure shows the hierarchical structure of adversarial attacks in the physical world, including image classification, object detection, and segmentation, with corresponding attack and defense strategies.

3.3 Case Studies of Adversarial Attacks in the Physical World

The physical world provides a unique set of challenges and opportunities for adversarial attacks, and various case studies have been conducted to demonstrate the potential risks and consequences of these attacks. According to [40], the Robust Physical Perturbations (RP2) method can generate robust physical perturbations that can consistently cause misclassification in a DNN-based classifier under a range of dynamic physical conditions. This highlights the potential risks and consequences of adversarial attacks in the physical world, where attackers can manipulate the physical environment to deceive computer vision systems [56]. The case studies of adversarial attacks in the physical world, as discussed in [57], include the proposed method of PadvFace, a robust physical attack framework that generates wearable adversarial stickers to deceive state-of-the-art face recognition models. Additionally, [43] presents a case study of designing an adversarial T-shirt to evade detection of moving persons, and the evaluation setting of the experiments, including a dataset of videos captured in different scenarios. The case studies of adversarial attacks in the physical world, as discussed in [32], highlight the potential risks and consequences of these attacks on object detectors, including the ability to manipulate the behavior of these systems and cause misclassification. By examining these case studies, researchers and practitioners can better appreciate the importance of developing secure and robust computer vision systems that can withstand adversarial attacks in the physical world [56, 43, 32].

4 Defenses Against Adversarial Attacks

4.1 Introduction to Defense Strategies

The development of effective defense strategies against adversarial attacks is crucial in ensuring the security and reliability of computer vision systems, as highlighted in [56], [57], and [7]. Defense strategies against adversarial attacks include adversarial training methods and input preprocessing and detection methods, which are essential in mitigating the potential risks and consequences of adversarial attacks [38], [22]. Understanding adversarial examples is crucial in developing effective defense strategies against adversarial attacks, as emphasized in [22]. The importance of understanding adversarial examples in the context of real-world applications is also highlighted in [22], [58], and [3].

As [4] discusses, the introduction to defense strategies against adversarial attacks is crucial in understanding the importance of developing robust defense strategies to protect against these attacks. The development of effective defense strategies against adversarial attacks is essential in mitigating the potential risks and consequences of these attacks, and continued research in this area is necessary to ensure the security and reliability of computer vision systems, as noted in [56], [57], and [4].

4.2 Adversarial Training Methods

Adversarial training methods are a crucial defense strategy against adversarial attacks, as they can improve the robustness of machine learning models and develop more effective defense methods [8], [20], [9], [10], [1], [2]. According to [8], adversarial training methods involve training a model on a dataset that includes adversarial examples, which can help the model learn to recognize and resist these types of attacks. The effectiveness of adversarial training methods has been demonstrated in various studies, including [20], which shows that these methods can improve the robustness of machine learning models against adversarial attacks by generating adversarial examples.

Furthermore, [9] discusses the use of adversarial training methods as a defense strategy against adversarial attacks, highlighting their strengths and limitations. Adversarial training methods have also been explored in [10], which demonstrate their potential for improving the robustness of object detection models against adversarial examples. Additionally, [1] and [2] provide a comprehensive overview of the strengths, weaknesses, and applications of adversarial training methods. The use of adversarial training methods as a defense strategy against adversarial attacks is also discussed in [20], which highlights the effectiveness of using methods for generating adversarial examples to improve the robustness of machine learning models.

4.3 Input Preprocessing and Detection Methods

Input preprocessing and detection methods are a crucial defense strategy against adversarial attacks, as they can help detect and mitigate the effects of these attacks on computer vision systems. According to [31], input preprocessing and detection methods can be used to detect adversarial examples by employing statistical hypothesis tests. This approach can help identify potential adversarial examples and prevent them from being processed by the computer vision system. Furthermore, [36] discusses the use of self-supervision at deployment time to defend against adversarial attacks, highlighting the potential of input preprocessing and detection methods in improving the robustness of computer vision systems.

Additionally, [18] proposes a method for detecting adversarial examples by comparing a model's predictions on feature-squeezed inputs with its predictions on the original inputs. This approach can help identify potential adversarial examples and prevent them from being processed by the computer vision system. The use of input preprocessing and detection methods as a defense strategy against adversarial attacks is essential in ensuring the security and reliability of computer vision systems, and continued research in this area is necessary to develop more effective defense strategies, as noted in [31], [36], and [18].

4.4 Challenges and Limitations of Defense Strategies

The development of effective defense strategies against adversarial attacks is crucial in ensuring the security and reliability of computer vision systems. However, as discussed in [37], the limitation of current methods is that they only evaluate the robustness of adversarial examples in a specific setup and do not provide a general solution to defend against them. Furthermore, [39] highlights the challenges and limitations of defense strategies against adversarial attacks, including the lack of a universal methodology for evaluating robustness. The survey in [47] also discusses the limitations of current methods for addressing natural adversarial examples, including their lack of effectiveness in mitigating their impact.

Additionally, [59] identifies gaps in current research, including the need for more robust and secure learning algorithms, more effective countermeasures, and a deeper understanding of the security properties of machine learning systems. The approach in [60] is also limited by its simplifying assumption, which may not generalize to all types of corruption or adversarial examples. As [33] emphasizes, defense strategies against adversarial attacks face challenges and limitations, including high computational cost, difficulty in solving the inner maximization problem, and a trade-off between robustness and generalization. The survey in [61] discusses the limitations of current defense methods, including the trade-off between performance on legitimate and adversarial examples.

Moreover, [62] highlights the weaknesses or limitations of current research, including the lack of a clear definition of robustness and the difficulty in evaluating the robustness of deep learning models. The challenges and limitations of defense strategies against adversarial attacks are also discussed

in [34], which includes the high codimension of the data manifold and the requirement of a large amount of computational resources to compute the decision axis. Furthermore, [25] notes that one limitation of DA-Attack is that it may not perform well on models with highly non-linear decision boundaries. The survey [14] discusses the weaknesses and limitations of current research, such as the lack of consideration of multiple data properties simultaneously and the need for more realistic problem setups. Additionally, [63] highlights the challenges and limitations of defense strategies, including nonlinear constraints and the need for robust defense mechanisms.

By understanding the challenges and limitations of defense strategies against adversarial attacks, researchers and practitioners can better appreciate the importance of developing robust and effective defense strategies to protect against these attacks. The development of effective defense strategies against adversarial attacks is essential in ensuring the security and reliability of computer vision systems, and continued research in this area is necessary to address the challenges and limitations of current methods, as noted in [37], [39], [47], [59], [60], [33], [61], [62], [34], [25], [14], [63].

5 Case Studies and Applications

This section provides an overview of various case studies and applications where visual adversarial attacks and defenses have significant implications, as highlighted in [64, 65, 66, 67, 68, 58, 8, 20, 9, 1]. The potential risks and consequences of adversarial attacks on various applications, including autonomous vehicles, surveillance systems, facial recognition technology, and other real-world scenarios, are a significant concern, as emphasized in [57, 43, 44, 4, 2, 1]. Understanding the potential risks and consequences of adversarial attacks on these applications is crucial in developing effective defense strategies to protect against these attacks, as noted in [8, 20, 9, 1].

5.1 Autonomous Vehicles

The potential risks and consequences of adversarial attacks on autonomous vehicles are significant, as autonomous vehicles rely heavily on computer vision systems to navigate and make decisions, and adversarial attacks on these systems can have severe consequences, including causing accidents or disrupting traffic flow, as highlighted in [64, 65, 66, 67, 68, 58, 8, 20, 9, 1]. The proposed algorithm, RP2, can be used to attack image classifiers in real-world scenarios such as autonomous vehicles, as discussed in [40], and the ABBA attack, as proposed in [45], can be used to attack autonomous vehicles, which use DNNs for image classification and object detection, and can cause potential risks and consequences. Furthermore, [65, 58] emphasizes the importance of understanding and addressing visual adversarial attacks and defenses in autonomous vehicles, highlighting the potential risks and consequences of these attacks.

5.2 Surveillance Systems

The potential risks and consequences of adversarial attacks on surveillance systems are significant, as surveillance systems rely heavily on computer vision systems to detect and track objects, and adversarial attacks on these systems can have severe consequences, including disrupting the functionality of the system or causing false alarms, as highlighted in [56, 43]. The proposed method of designing an adversarial T-shirt to evade detection of moving persons demonstrates the potential risks and consequences of adversarial attacks on surveillance systems, as discussed in [43], and [32] emphasizes the importance of robustness improvements for object detection models, which are commonly used in surveillance systems, to mitigate the potential risks and consequences of adversarial attacks.

5.3 Facial Recognition Technology

The potential risks and consequences of adversarial attacks on facial recognition technology are significant, as facial recognition technology is widely used in various applications, including security, surveillance, and authentication, and adversarial attacks on these systems can have severe consequences, including dodging and impersonation attacks, as highlighted in [57, 44]. The proposed method in [57] demonstrates the potential risks and consequences of adversarial attacks on facial recognition technology, and [44] emphasizes the potential risks and consequences of adversarial

attacks on facial recognition technology in the black-box setting, where the attacker has limited knowledge of the system’s architecture and parameters.

5.4 Other Real-World Scenarios

Visual adversarial attacks and defenses are relevant in various real-world scenarios beyond autonomous vehicles, surveillance systems, and facial recognition technology, as emphasized in [4, 2, 1]. For instance, in the field of healthcare, medical imaging systems can be vulnerable to adversarial attacks, which can lead to misdiagnosis or incorrect treatment, as discussed in [2], and in the field of finance, adversarial attacks on image classification systems can be used to manipulate financial transactions or steal sensitive information, as highlighted in [1]. The use of adversarial training methods, as discussed in [8], can be effective in improving the robustness of machine learning models against adversarial attacks in these scenarios, and [20] highlights the importance of input preprocessing and detection methods in mitigating the effects of adversarial attacks in real-world scenarios. By understanding the relevance of visual adversarial attacks and defenses in various real-world scenarios, researchers and practitioners can better appreciate the importance of developing robust defense strategies to protect against these attacks and ensure the security and reliability of various applications, as noted in [9, 20, 1].

6 Conclusion

6.1 Future Directions and Open Research Questions

The realm of visual adversarial attacks and defenses is continually evolving, with a plethora of future directions and open research questions awaiting exploration, as highlighted by the works of [20, 9, 1, 2, 3]. A crucial aspect of this evolution involves the development of more robust defense strategies, which can be achieved through the investigation of novel architectures and techniques to mitigate the impact of adversarial examples, as discussed in [1, 2]. Furthermore, exploring new types of attacks and examining the potential applications in real-world scenarios, such as autonomous vehicles, surveillance systems, and facial recognition technology, are essential for advancing the field, as noted in [9, 1, 2]. The application of visual adversarial attacks and defenses in other areas of machine learning, including natural language processing and reinforcement learning, presents a promising direction for future research, as suggested by [69]. By delving into these future directions and open research questions, researchers can gain a deeper understanding of the complexities inherent to visual adversarial attacks and defenses, ultimately leading to the development of more effective defense strategies, as emphasized by [3, 69, 33].

The potential applications and challenges of visual adversarial attacks and defenses in various fields necessitate further exploration, with a focus on developing more effective defense strategies and examining the utility of these strategies in other applications, as discussed in [9, 1, 2]. Future research endeavors should prioritize improving the efficiency and effectiveness of defense strategies, as well as exploring the application of these strategies in other domains, as noted in [3]. The development of more robust and efficient defense mechanisms, the exploration of novel attack strategies, and the application of adversarial machine learning to other areas of machine learning are also vital, as highlighted by [69, 33]. Additionally, investigating new methods beyond adversarial training for adversarial robustness and achieving desirable robustness and accuracy in deep learning models are critical for future research, as emphasized by [33].

6.2 Future Directions and Applications

The field of visual adversarial attacks and defenses holds significant potential for growth and development, with numerous potential future directions and applications awaiting exploration, as discussed in [69, 33]. Future research directions include the development of more robust and efficient defense mechanisms, the exploration of new attack strategies, and the application of adversarial machine learning to other areas of machine learning, as noted in [69]. The potential applications of visual adversarial attacks and defenses in various fields, such as autonomous vehicles, surveillance systems, and facial recognition technology, necessitate further examination, as highlighted by [9, 1, 2]. The application of visual adversarial attacks and defenses in other areas of machine learning, including natural language processing and reinforcement learning, presents a promising direction for future

research, as suggested by [69]. By exploring these future directions and applications, researchers can gain a deeper understanding of the complexities inherent to visual adversarial attacks and defenses, ultimately leading to the development of more effective defense strategies to protect against these attacks, as emphasized by [3, 69, 33].

www.SurveyX.cn

References

- [1] Shuangtao Li, Lin Bai, Yanlin Peng, and Yuanke Chen. Learning more robust features with adversarial training. *arXiv preprint arXiv:1804.07757*, 2018.
- [2] Dou Goodman. Transferability of adversarial examples to attack cloud-based image classifier service. *arXiv preprint arXiv:2001.03460*, 2020.
- [3] Zhe Zhou, Di Tang, Kehuan Zhang, Weili Han, Xiangyu Liu, and XiaoFeng Wang. Invisible mask: Practical attacks on face recognition with infrared. *arXiv preprint arXiv:1803.04683*, 2018.
- [4] Zhe Zhao, Fu Song, Guangke Chen, Jingyi Wang, Jun Sun, Taishan Li, and Tong Liu. Attack as detection: Using adversarial attack methods to detect abnormal examples. *ACM Transactions on Software Engineering and Methodology*, 2023.
- [5] Xiaosen Wang, Chuanbiao Song, John E. Hopcroft, Kun He, and Liwei Wang. At-gan: An adversarial generator model for non-constrained adversarial examples. *arXiv preprint arXiv:1904.07793*, 2019.
- [6] Andrew Ilyas, Anish Athalye, Jessy Lin, and Logan Engstrom. Black-box adversarial attacks with limited queries and information. *arXiv preprint arXiv:1804.08598*, 2018.
- [7] Nupur Kumari, Abhishek Sinha, Balaji Krishnamurthy, Harshitha Machiraju, Mayank Singh, and Vineeth N Balasubramanian. Harnessing the vulnerability of latent layers in adversarially trained models. 2019.
- [8] Shuo Wang, Alsharif Abuadbba, Carsten Rudolph, Marthie Grobler, and Surya Nepal. Adversarial detection by latent style transformations. *IEEE Transactions on Information Forensics and Security*, 2022.
- [9] Yunhan Jia, Junjie Shen, Qi Alfred Chen, Tao Wei, Yantao Lu, and Zhenyu Zhong. Fooling detection alone is not enough: First adversarial attack against multiple object tracking. *arXiv preprint arXiv:1905.11026*, 2019.
- [10] Xiaofeng Lu, Andrzej Stefan Ceglowski, Chuan Qin, Guangling Sun, XU Wen-bo, and Yuying Su. Complete defense framework to protect deep neural networks against adversarial examples. *Mathematical Problems in Engineering*, 2020.
- [11] Anirban Chakraborty, Anupam Chattopadhyay, Debdeep Mukhopadhyay, Manaar Alam, and Vishal Dey. Adversarial attacks and defences: A survey. *arXiv preprint arXiv:1810.00069*, 2018.
- [12] Liam Fowl, Jonas Geiping, Micah Goldblum, Ping yeh Chiang, Tom Goldstein, and Wojtek Czaja. Adversarial examples make strong poisons. *arXiv preprint arXiv:2106.10807*, 2021.
- [13] Shuo Wang, Carsten Rudolph, Marthie Grobler, Shangyu Chen, Tianle Chen, and Surya Nepal. Generating semantic adversarial examples via feature manipulation. *arXiv preprint arXiv:2001.02297*, 2020.
- [14] Peiyu Xiong, Jaskeerat Singh Sarin, Julia Rubin, Michael Tegegn, and Shubhraneel Pal. It is all about data: A survey on the effects of data on adversarial robustness. *ACM Computing Surveys*, 2023.
- [15] Yinpeng Dong, Baoyuan Wu, Hang Su, Jun Zhu, Tong Zhang, Wei Liu, and Zhifeng Li. Efficient decision-based black-box adversarial attacks on face recognition. 2019.
- [16] Kazuya Kakizaki and Kosuke Yoshida. Adversarial image translation: Unrestricted adversarial examples in face recognition systems. *arXiv preprint arXiv:1905.03421*, 2019.
- [17] Lu Chen and Wei Xu. Attacking optical character recognition (ocr) systems with adversarial watermarks. *arXiv preprint arXiv:2002.03095*, 2020.
- [18] Weilin Xu, David Evans, and Yanjun Qi. Feature squeezing: Detecting adversarial examples in deep neural networks. 2018.

-
- [19] Kejiang Chen, Hang Zhou, Hui Xue, Nenghai Yu, Weiming Zhang, Xiaofeng Mao, Yuan He, Yuefeng Chen, and Yuhong Li. Self-supervised adversarial training. 2020.
- [20] Florian Tramèr, Dan Boneh, Giancarlo Pellegrino, Gili Rusak, and Pascal Dupré. Adversarial: Perceptual ad blocking meets adversarial machine learning. 2019.
- [21] Anish Athalye and Nicholas Carlini. On the robustness of the cvpr 2018 white-box adversarial example defenses. *arXiv preprint arXiv:1804.03286*, 2018.
- [22] Jiajun Lu, Evan Fabry, and Hussein Sibai. Adversarial examples that fool detectors. *arXiv preprint arXiv:1712.02494*, 2017.
- [23] Maung Maung, April Pyone, and Hitoshi Kiya. Encryption inspired adversarial defense for visual classification. 2020.
- [24] Jon Vellido, José A. Lozano, and Roberto Santana. Adversarial attacks in explainable machine learning: A survey of threats against models and humans. *Wiley Interdisciplinary Reviews Data Mining and Knowledge Discovery*, 2024.
- [25] Tianjin Huang, Mykola Pechenizkiy, Vlado Menkovski, Yuhao Wang, and Yulong Pei. Direction-aggregated attack for transferable adversarial examples. *ACM Journal on Emerging Technologies in Computing Systems*, 2022.
- [26] Xintao Ding, Prosanta Gope, Qingde Li, Yonglong Luo, and Yongqiang Cheng. Consensus adversarial defense method based on augmented examples. *IEEE Transactions on Industrial Informatics*, 2022.
- [27] Zhaoxia Yin, Abel Andrew, Li Chen, and Shaowei Zhu. Reversible attack based on local visible adversarial perturbation. *Multimedia Tools and Applications*, 2023.
- [28] Chia-Yi Hsu, Chia-Mu Yu, Pin-Yu Chen, Sijia Liu, and Songtao Lu. Adversarial examples can be effective data augmentation for unsupervised machine learning. *Proceedings of the AAAI Conference on Artificial Intelligence*, 2022.
- [29] Ke Sun, Zhanxing Zhu, and Zhouchen Lin. Towards understanding adversarial examples systematically: Exploring data size, task and model factors. *arXiv preprint arXiv:1902.11019*, 2019.
- [30] Simant Dube. High dimensional spaces, deep learning and adversarial examples. *arXiv preprint arXiv:1801.00634*, 2018.
- [31] Yang Song, Nate Kushman, Sebastian Nowozin, Stefano Ermon, and Taesup Kim. Pixeldefend: Leveraging generative models to understand and defend against adversarial examples. *arXiv preprint arXiv:1710.10766*, 2017.
- [32] Yusheng Zhao, Huanqian Yan, and Xingxing Wei. Object hider: Adversarial patch attack against object detectors. *arXiv preprint arXiv:2010.14974*, 2020.
- [33] Tao Bai, Bihan Wen, Jinqi Luo, Jun Zhao, and Qian Wang. Recent advances in adversarial training for adversarial robustness. *arXiv preprint arXiv:2102.01356*, 2021.
- [34] Marc Khoury and Dylan Hadfield-Menell. On the geometry of adversarial examples. *arXiv preprint arXiv:1811.00525*, 2018.
- [35] Ranjie Duan, A. K. Qin, Dantong Niu, Yuan He, Yuefeng Chen, and Yun Yang. Advdrop: Adversarial attack to dnns by dropping information. *2021 IEEE/CVF International Conference on Computer Vision (ICCV)*, 2021.
- [36] Chengzhi Mao, Carl Vondrick, Hao Wang, Junfeng Yang, and Mia Chiquier. Adversarial attacks are reversible with natural supervision. *2021 IEEE/CVF International Conference on Computer Vision (ICCV)*, 2021.
- [37] Jiajun Lu, David Forsyth, Evan Fabry, and Hussein Sibai. No need to worry about adversarial examples in object detection in autonomous vehicles. *arXiv preprint arXiv:1707.03501*, 2017.

-
- [38] Rey Reza Wiyatno, Anqi Xu, Archy O. de Berker, and Ousmane Dia. Adversarial examples in modern machine learning: A review. *arXiv preprint arXiv:1911.05268*, 2019.
- [39] Xiaoyong Yuan, Pan He, Qile Zhu, and Xiaolin Li. Adversarial examples: Attacks and defenses for deep learning. *arXiv preprint arXiv:1712.07107*, 2017.
- [40] Ivan Evtimov, Amir Rahmati, Atul Prakash, Bo Li, Dawn Song, Earlence Fernandes, Kevin Eykholt, Dawn Song, and Tadayoshi Kohno. Robust physical-world attacks on deep learning models. *arXiv preprint arXiv:1707.08945*, 2017.
- [41] Anish Athalye, Andrew Ilyas, Kevin Kwok, and Logan Engstrom. Synthesizing robust adversarial examples. *arXiv preprint arXiv:1707.07397*, 2017.
- [42] Chaowei Xiao, Bo Li, Dawn Song, Jun-Yan Zhu, Mingyan Liu, and Warren He. Spatially transformed adversarial examples. *arXiv preprint arXiv:1801.02612*, 2018.
- [43] Kaidi Xu, Gaoyuan Zhang, Hongge Chen, Mengshu Sun, Pin-Yu Chen, Quanfu Fan, Sijia Liu, Xue Lin, and Yanzhi Wang. Adversarial t-shirt! evading person detectors in a physical world. *arXiv preprint arXiv:1910.11099*, 2019.
- [44] Yigit Alparslan, Rachel Greenstadt, Jeremy Keim-Shenk, Rachel Greenstadt, and Shweta Khade. Adversarial attacks on convolutional neural networks in facial recognition domain. *arXiv preprint arXiv:2001.11137*, 2020.
- [45] Qing Guo, Bing Yu, Felix Juefei-Xu, Jian Wang, Lei Ma, Wei Feng, Xiaofei Xie, and Yang Liu. Watch out! motion is blurring the vision of your deep neural networks. *arXiv preprint arXiv:2002.03500*, 2020.
- [46] Bo Yang, Hengwei Zhang, Jindong Wang, Kaiyong Xu, Yuchen Zhang, and Zheming Li. Adversarial example generation with adabelief optimizer and crop invariance. *Applied Intelligence*, 2022.
- [47] Aníbal Pedraza, Gloria Bueno, and Óscar Déniz. Really natural adversarial examples. *International Journal of Machine Learning and Cybernetics*, 2021.
- [48] Nicolas Papernot, Abhibhav Garg, Alexander Matyasko, Alexey Kurakin, Aurko Roy, Cihang Xie, David Berthelot, Fartash Faghri, Ian Goodfellow, Jonas Rauber, Jonathan Uesato, Karen Hambardzumyan, Nicholas Carlini, Patrick McDaniel, Paul Hendricks, Reuben Feinman, Rujun Long, Ryan Sheatsley, T. B. Brown, Vahid Behzadan, Willi Gierke, Yash Sharma, Yi-Lin Juang, Yinpeng Dong, Zhi Li, and Zhishuai Zhang. Technical report on the cleverhans v2.1.0 adversarial examples library. *arXiv preprint arXiv:1610.00768*, 2016.
- [49] Chang Song, Chunpeng Wu, Hai Li, Hsin-Pai Cheng, Huanrui Yang, Qing Wu, Sicheng Li, and Yiran Chen. Mat: A multi-strength adversarial training method to mitigate adversarial attacks. 2018.
- [50] Zifei Zhang, Bin Yan, Jian Chen, Kai Qiao, Lingyun Jiang, and Linyuan Wang. Advjnd: Generating adversarial examples with just noticeable difference. *Lecture notes in computer science*, 2020.
- [51] Ian Goodfellow. Defense against the dark arts: An overview of adversarial example security research and future research directions. *arXiv preprint arXiv:1806.04169*, 2018.
- [52] Yuchen Zhang and Percy Liang. Defending against whitebox adversarial attacks via randomized discretization. *arXiv preprint arXiv:1903.10586*, 2019.
- [53] Qilong Zhang, Hui Xue, Jingkuan Song, Lianli Gao, Xiaodan Li, Yuan He, and Yuefeng Chen. Beyond imagenet attack: Towards crafting adversarial examples for black-box domains. *arXiv preprint arXiv:2201.11528*, 2022.
- [54] Shasha Li, Ajaya Neupane, Amit K. Roy Chowdhury, Ananthram Swami, Chengyu Song, Srikanth V. Krishnamurthy, and Sujoy Paul. Adversarial perturbations against real-time video classification systems. 2019.

-
- [55] Anibal Pedraza, Gloria Bueno, and Óscar Déniz. Approaching adversarial example classification with chaos theory. *Entropy*, 2020.
- [56] Jiakai Wang, Aishan Liu, Shiyu Tang, Shunchang Liu, Xianglong Liu, and Zixin Yin. Dual attention suppression attack: Generate adversarial camouflage in physical world. 2021.
- [57] Baoyuan Wu, Jue Wang, Shirui Pan, Xin Zheng, Yanbo Fan, and Yong Zhang. Robust physical-world attacks on face recognition. *Pattern Recognition*, 2022.
- [58] Zhe Zhao, Fu Song, Guangke Chen, Jingyi Wang, Jun Sun, and Yiwei Yang. Attack as defense: characterizing adversarial examples using robustness. 2021.
- [59] Battista Biggio and Fabio Roli. Wild patterns: Ten years after the rise of adversarial machine learning. *Pattern Recognition*, 2018.
- [60] Nic Ford, Dogus Cubuk, Justin Gilmer, and Nicolas Carlini. Adversarial examples are a natural consequence of test error in noise. *arXiv preprint arXiv:1901.10513*, 2019.
- [61] Junhao Dong, Hao Chen, Jianhuang Lai, Junxi Chen, and Xiaohua Xie. Survey on adversarial attack and defense for medical image analysis: Methods and challenges. *ACM Computing Surveys*, 2024.
- [62] Yaochu Jin and Jia Lin Liu. A comprehensive survey of robust deep learning in computer vision. *Journal of Automation and Intelligence*, 2023.
- [63] Jiangnan Li, Hairong Qi, Jinyuan Sun, Kevin Tomsovic, and Yingyuan Yang. Conaml: Constrained adversarial machine learning for cyber-physical systems. 2021.
- [64] Jiayang Liu, Jun Sakuma, Kazuto Fukuchi, Weiming Zhang, and Youhei Akimoto. Unauthorized ai cannot recognize me: Reversible adversarial example. *Pattern Recognition*, 2022.
- [65] Abderrahmen Amich and Birhanu Eshete. Morphence: Moving target defense against adversarial examples. *Annual Computer Security Applications Conference*, 2021.
- [66] Sanjit A. Seshia, Somesh Jha, and Tommaso Dreossi. Semantic adversarial deep learning. *IEEE Design and Test*, 2020.
- [67] AprilPyone MaungMaung, Hitoshi Kiya, and Yuma Kinoshita. Adversarial robustness by one bit double quantization for visual classification. *IEEE Access*, 2019.
- [68] Andras Rozsa and Terrance E. Boult. Improved adversarial robustness by reducing open space risk via tent activations. *arXiv preprint arXiv:1908.02435*, 2019.
- [69] Guofu Li, Cao Ning, Jin Li, Pengjia Zhu, Zheming Yang, and Zhiyi Chen. Security matters: A survey on adversarial machine learning. *arXiv preprint arXiv:1810.07339*, 2018.

Disclaimer:

SurveyX is an AI-powered system designed to automate the generation of surveys. While it aims to produce high-quality, coherent, and comprehensive surveys with accurate citations, the final output is derived from the AI's synthesis of pre-processed materials, which may contain limitations or inaccuracies. As such, the generated content should not be used for academic publication or formal submissions and must be independently reviewed and verified. The developers of SurveyX do not assume responsibility for any errors or consequences arising from the use of the generated surveys.

www.SurveyX.cn